

CS-1217 Operating Systems

Spring 2024

Lab 1

March 4, 2024

Exercise 1

```
# Jump to next instruction, but in 32-bit code segment.
# Switches processor into 32-bit mode.
ljmp .. $PROT_MODE_CSEG, $protcseg
```

1. The instruction `ljmp .. $PROT_MODE_CSEG, $protcseg` in *boot.S* causes the switch from 16-bit to 32-bit mode.

2. For the last instruction of the boot loader executed,

- in *main.c*:

```
// call the entry point from the ELF header
// note: does not return!
((void (*)(void)) (ELFHDR->e_entry))();
```

- in *boot.asm*:

```
((void (*)(void)) (ELFHDR->e_entry))();
7d6b: ff 15 18 00 01 00    call    *0x10018
```

The first instruction to be loaded is:

```
f010000c: 66 c7 05 72 04 00 00    movw    $0x1234,0x472
```

3. Since the last instruction the boot loader executes is `call *0x10018`, the first instruction of the kernel should be at this address. When we look at this address using `gdb`, we get:

```
(gdb) x/1x 0x10018
0x10018:    0x0010000c
```

Hence, the first instruction is at `0x0010000c`

4. The boot loader initializes pointers `ph` and `eph` to point to the program headers in the ELF header. `ph` is set to point to the first program header by calculating its address relative to the start of the ELF header while `eph` is set to point to the end of the program headers by adding the number of program headers (`e_phnum`) to `ph`.

```
// load each program segment (ignores ph flags)
ph = (struct Proghdr *) ((uint8_t *) ELFHDR + ELFHDR->e_phoff);
eph = ph + ELFHDR->e_phnum;
for (; ph < eph; ph++)
    // p_pa is the load address of this segment (as well
    // as the physical address)
    readseg(ph->p_pa, ph->p_memsz, ph->p_offset);
```

Exercise 2

1. *console.c* exports cputchar, getchar, and iscons, while cputchar is used as a parameter when *printf.c* calls vprintfmt from *printfmt.c*
2. In the *console.c* file, it verifies if the cursor has reached the buffer's end, i.e. the screen is full. If it has, the buffer is scrolled up by one row, the last row is cleared with spaces, and the cursor is positioned at the start of the last row to make space for newer information.
3.
 - In the call to cprintf(), *fmt* points to the format string of its arguments while *ap* points to the variable arguments after *fmt*.
 - For this part, I modified the *monitor.c* and added the snippet to it, then ran gdb and got:

```

cprintf (fmt=0xf0101ad2 "x %d, y %x, z %d\n")
vcprintf (fmt=0xf0101ad2 "x %d, y %x, z %d\n", ap=0xf0115f64 "\001")
cons_putc (c=120)
cons_putc (c=32)
va_arg(*ap, int)
Hardware watchpoint 4: ap
Old value = 0xf0115f64 "\001"
New value = 0xf0115f68 "\003"
cons_putc (c=49)
cons_putc (c=44)
cons_putc (c=32)
cons_putc (c=121)
cons_putc (c=32)
va_arg(*ap, int)
Hardware watchpoint 4: ap
Old value = 0xf0115f68 "\003"
New value = 0xf0115f6c "\004"
cons_putc (c=51)
cons_putc (c=44)
cons_putc (c=32)
cons_putc (c=122)
cons_putc (c=32)
va_arg(*ap, int)
Hardware watchpoint 4: ap
Old value = 0xf0115f6c "\004"

```

```
New value = 0xf0115f70 "T\034\020?\214_\021??\027\020??_\021??\027\020?_\021?_\021?"
cons_putc (c=52)
cons_putc (c=10)
```

4. The output is He110 World, because 57616=0xe110, so the first half of output is He110 because 57616 is read in hexadecimal. i=0x00646c72 is treated as a string, so it will be printed as 'r'=(char)0x72, 'l'=(char)0x6c, 'd'=(char)0x64, and 0x00 is treated as a mark of end of string.

We will see He110, Wo in a big-endian machine. 57616 will still be read as e110 because only its numeric value matters when being printed. However, when i = 0x00646c72 is treated as the string, the 0x00 terminates the string at Wo

5. After 'y=', the decimal value of 4 bytes right above where 3 is placed in the stack will be printed.
6. GCC changed its calling convention so that it pushed arguments on the stack in declaration order, so that the last argument is pushed last. We can just store the number of arguments as an integer. To adapt cprintf for GCC's new calling convention, its interface would need to include the number of arguments being passed by modifying cprintf to accept a count of arguments along with a format string and a variable argument list. The caller would need to provide the count of the arguments while invoking the function. For example, the interface of cprintf could be updated to something like cprintf(const char *fmt, int num_args, ...).

Exercise 3

Using *kernel.asm*, we can notice that the starting address for the `test_backtrace` is `0xf0100040`.

Using `gdb`, we get:

```
+ symbol-file kernel
(gdb) b *0xf0100040
Breakpoint 1 at 0xf0100040: file kern/init.c, line 13.
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0xf0100040 <test_backtrace>: push    %ebp

Breakpoint 1, test_backtrace (x=5) at kern/init.c:13
13      {
(gdb) i r
eax                0x0          0
ecx                0x3d4        980
edx                0x3d5        981
ebx                0xf0111308    -267316472
esp                0xf010ffdc    0xf010ffdc
ebp                0xf010fff8    0xf010fff8
esi                0x10094      65684
edi                0x0          0
eip                0xf0100040    0xf0100040 <test_backtrace>
eflags             0x46         [ PF ZF ]
cs                 0x8          8
---Type <return> to continue, or q <return> to quit---
ss                 0x10         16
ds                 0x10         16
es                 0x10         16
fs                 0x10         16
gs                 0x10         16
(gdb) c
Continuing.
=> 0xf0100040 <test_backtrace>: push    %ebp
```

```

(gdb) c
Continuing.
=> 0xf0100040 <test_backtrace>: push  %ebp

Breakpoint 1, test_backtrace (x=4) at kern/init.c:13
13  {
(gdb) i r
eax          0x4      4
ecx          0x3d4    980
edx          0x3d5    981
ebx          0xf0111308 -267316472
esp          0xf010ffbc 0xf010ffbc
ebp          0xf010ffd8 0xf010ffd8
esi          0x5      5
edi          0x0      0
eip          0xf0100040 0xf0100040 <test_backtrace>
eflags      0x92     [ AF SF ]
cs          0x8      8
---Type <return> to continue, or q <return> to quit---
ss          0x10     16
ds          0x10     16
es          0x10     16
fs          0x10     16
gs          0x10     16
(gdb) █

```

The difference of ebp between the two breakpoints is 0x20, so every time it pushes 8 4-byte words as follows:

```

return address
saved ebp
saved ebx
abandoned
abandoned
abandoned
abandoned
var x for calling next test_backtrace

```

The return instruction pointer typically points to the instruction after the call instruction (why?) → The return instruction pointer points to the instruction after the call instruction because it ensures sequential execution, allowing the CPU to resume executing code sequentially after a function call, and it aligns with function call conventions, where the return address is pushed onto the stack before jumping to the called function, simplifying the management of function calls and returns for efficient control flow.

Why can't the backtrace code detect how many arguments there actually are? How could this limitation be fixed? → The backtrace code cannot detect the number of arguments passed to a

function because C does not provide built-in mechanisms to check the number of arguments at runtime. The number of arguments passed to a function is fixed and determined at compile time, making it impossible for the backtrace code to dynamically determine the number of arguments. We could use a different calling convention or function signature that includes information about the number of arguments. You could define a struct to hold both the number of arguments and the arguments themselves, passing this struct to the backtrace function, giving it the number of arguments and their values.